

QRNG comparison report

stagea_ibm_marrakesh_20260714-210803_processed.txt vs optimal_harvested_ibm_marrakesh_20260714-215818_
Generated 2026-07-15 09:17


Bits used per file: ALL | NIST bits/partition: 1,000,000 | ML windows [8, 16, 24], 8 folds | Markov orders [1, 2, 4]

stagea_ibm_marrakesh_20260714-210803_processed.txt	stagea_ibm_marrakesh_20260714-210803_processed.txt	optimal_harvested_ibm_marrakesh_20260714-215818_	optimal_harvested_ibm_marrakesh_20260714-215818_	optimal_harvested_ibm_marrakesh_20260714-215818_
Total bits	10,140,000	7,605,000	20,000	15,000
Global bias p1-0.5	0.0073	0.0002	0.0050	0.0031
Min-entropy / bit	0.9790	0.9994	0.9856	0.9910
NIST pass rate	0.48	1.00	1.00	1.00
NIST passed / run	29/60	66/66	45/45	45/45
Next-bit verdict	PASS	PASS	PASS	PASS
Next-bit worst AUC	0.5000	0.5000	0.5000	0.5000
Markov verdict	PASS	PASS	PASS	PASS
Markov max excess	0.0007	0.0000	0.0078	0.0000

All four streams side by side: each input file contributes its RAW stream and its WHITENED (SHA-512, 75% kept) stream. Lower bias is better; min-entropy near 1.0 is better;
NIST pass rate near 1.0 is better; next-bit / Markov verdict PASS means no detectable structure (accuracy ~0.5 with p >= 0.01 and a 95% CI that includes 0.5).

Bias per segment (20 splits)

Local balance along each stream — all four streams overlaid



Each stream is cut into 20 equal segments; the plotted value is how far that segment's fraction of ones sits from 0.5. Solid lines = raw, dashed = whitened. A good source is LOW and FLAT: no segment drifts and there is no rising/falling trend. Whitened lines should sit near the bottom regardless of how their raw counterpart behaves — that is the extractor removing residual imbalance. Reading tip: compare the two SOLID (raw) lines to judge the hardware; the dashed lines show what post-processing recovers.

Bias per segment — data ($|p1 - 0.5|$ per split)


Data table (same numbers as the chart)

Split	sea_marrakesh_20260714-215	sea_marrakesh_20260714-215	sea_marrakesh_20260714-215	sea_marrakesh_20260714-215
1	0.0061	0.0008	0.0150	0.0000
2	0.0063	0.0009	0.0060	0.0187
3	0.0077	0.0009	0.0070	0.0000
4	0.0072	0.0001	0.0040	0.0200
5	0.0068	0.0008	0.0060	0.0133
6	0.0075	0.0005	0.0080	0.0120
7	0.0087	0.0003	0.0000	0.0133
8	0.0083	0.0012	0.0060	0.0053
9	0.0071	0.0007	0.0010	0.0107
10	0.0073	0.0010	0.0140	0.0187
11	0.0074	0.0014	0.0040	0.0093
12	0.0080	0.0004	0.0070	0.0280
13	0.0083	0.0005	0.0020	0.0067
14	0.0070	0.0005	0.0060	0.0280
15	0.0074	0.0005	0.0410	0.0227
16	0.0070	0.0003	0.0080	0.0067
17	0.0077	0.0014	0.0160	0.0360
18	0.0076	0.0009	0.0150	0.0147
19	0.0067	0.0006	0.0090	0.0080
20	0.0069	0.0007	0.0050	0.0147

Same values as the chart above. Lower is better; look for any column that is systematically larger (a more biased stream) or any single spiking row (localised imbalance).

Global bias & min-entropy — raw vs whitened

Whole-stream imbalance for all four streams



Whole-stream bias (one number per stream). The whitened bars should be far lower than their raw counterparts — direct evidence the SHA-512 extractor removes imbalance while keeping 75% of the data. Min-entropy per bit (table) converts bias to 'how many bits of true randomness per output bit'; 1.0 is ideal.

Global bias & min-entropy — data

Data table (same numbers as the chart)


Marrakesh_20260714-215	Marrakesh_20260714-215	Marrakesh_20260714-215	Marrakesh_20260714-215	Marrakesh_20260714-215
Global bias	0.0073	0.0002	0.0050	0.0031
Min-entropy/bit	0.9790	0.9994	0.9856	0.9910
Total bits	10,140,000	7,605,000	20,000	15,000

Same numbers as the chart. Compare each raw stream with its whitened version to quantify how much structure the extractor removed.

NIST SP 800-22 — partitions passed per test

Out of 5 partitions; all four streams grouped per test

Overall pass rate — stagea_marrakesh_20260714-210803_processed raw: 0.48 stagea_marrakesh_20260714-210803_processed white: 1.00 optim



The full NIST SP 800-22 battery runs on 5 partitions of each stream; a bar of 5 means the test passed in every partition ($p > 0.01$ each). Longer is better. Tests that need more bits than a partition provides are skipped by the library (so the test list can differ between streams if bit counts differ). Whitened streams should reach 5 across the board.

NIST SP 800-22 — data (partitions passed / run)


Data table (same numbers as the chart)

	NIST test				
	tea_marrakesh_20260714-215803	tea_marrakesh_20260714-215803	tea_marrakesh_20260714-215803	tea_marrakesh_20260714-215803	tea_marrakesh_20260714-215803
Approximate Entropy	0/5	5/5	5/5	5/5	5/5
Binary Matrix Rank	5/5	5/5	—	—	—
Cumulative Sums	0/5	5/5	5/5	5/5	5/5
Discrete Fourier Transform	5/5	5/5	5/5	5/5	5/5
Frequency Within Block	0/5	5/5	5/5	5/5	5/5
Linear Complexity	5/5	5/5	—	—	—
Longest Run Ones In A Block	4/5	5/5	5/5	5/5	5/5
Maurers Universal	5/5	5/5	—	—	—
Monobit	0/5	5/5	5/5	5/5	5/5
Non Overlapping Template Match	5/5	5/5	5/5	5/5	5/5
Random Excursion	—	3/3	—	—	—
Random Excursion Variant	—	3/3	—	—	—
Runs	0/5	5/5	5/5	5/5	5/5
Serial	0/5	5/5	5/5	5/5	5/5
OVERALL pass rate	0.48	1.00	1.00	1.00	1.00

Same values as the chart. '4/5' means the test passed in 4 of the 5 partitions. '—' means the test was not eligible for that stream's length.

Next-bit predictability — strict, bias-robust (AUC)

AUC vs window size; 0.5 = no learnable structure



A logistic-regression model predicts each bit from the previous W bits, 8-fold time-ordered CV, for W in $[8, 16, 24]$. We use ROC AUC as the primary signal because it is BIAS-ROBUST: a biased-but-independent stream still gives AUC ~ 0.5 , whereas plain accuracy would rise just from guessing the majority bit. Structure is flagged only if a window's AUC 95% CI lies ABOVE 0.5 (or accuracy beats the majority baseline). Verdict in the legend is the strict AND over all windows. (Raw bias itself is reported separately in the bias analysis.)

Next-bit predictability — data (per stream × window)


Data table (same numbers as the chart)

	Stream	Window	Accuracy	Majority base	AUC	AUC 95% CI	p	pass
rrakesh_20260714-2108		8	0.5052	0.5054	0.5000	[0.497,0.503]	0.000	PASS
rrakesh_20260714-2108		16	0.5052	0.5054	0.4982	[0.495,0.501]	0.000	PASS
rrakesh_20260714-2108		24	0.5016	0.5054	0.4966	[0.494,0.500]	0.113	PASS
rrakesh_20260714-21080		8	0.5004	0.5009	0.4995	[0.497,0.502]	0.391	PASS
rrakesh_20260714-21080		16	0.5003	0.5009	0.4992	[0.496,0.502]	0.421	PASS
rrakesh_20260714-21080		24	0.4985	0.5009	0.4977	[0.495,0.501]	0.870	PASS
ed_mrrakesh_20260714		8	0.4983	0.5050	0.4994	[0.488,0.511]	0.637	PASS
ed_mrrakesh_20260714		16	0.4965	0.5050	0.4961	[0.485,0.507]	0.761	PASS
ed_mrrakesh_20260714		24	0.4948	0.5050	0.4949	[0.484,0.506]	0.853	PASS
ed_mrrakesh_20260714-		8	0.4989	0.5031	0.4916	[0.479,0.505]	0.578	PASS
ed_mrrakesh_20260714-		16	0.4915	0.5031	0.4835	[0.470,0.497]	0.932	PASS
ed_mrrakesh_20260714-		24	0.4969	0.5031	0.4933	[0.480,0.506]	0.706	PASS

'Majority base' = accuracy obtainable by always guessing the more common bit (pure bias). A window PASSES when AUC's CI includes/sits at 0.5 AND accuracy does not significantly exceed the majority baseline.

Markov dependency — multi-order (bias-robust)

Accuracy ABOVE the majority baseline; 0 = no dependency



A k-th order Markov predictor guesses the next bit from the previous k bits. We plot how far its accuracy exceeds the MAJORITY (bias-only) baseline, so pure bias reads as ~0 — only genuine short-range DEPENDENCY lifts a bar above the green line. Whitened streams should sit at 0 at every order. (Overall bias is reported separately.)

Markov dependency — data (accuracy, baseline, excess per order)

Data table (same numbers as the chart)

	Stream	o1 acc	o2 acc	o4 acc	base	o1 excess	o2 excess	o4 excess	verdict
rakesh_20260714-2108		0.5069	0.5069	0.5069	0.5063	0.0007	0.0007	0.0007	PASS
rakesh_20260714-2108		0.4993	0.4998	0.5000	0.5004	-0.0011	-0.0006	-0.0004	PASS
ed_marrakesh_20260714		0.5083	0.5128	0.5003	0.5050	0.0033	0.0078	-0.0047	PASS
d_marrakesh_20260714		0.4980	0.5004	0.5006	0.5031	-0.0051	-0.0027	-0.0026	PASS

'base' = majority baseline (pure bias). 'excess' = accuracy – baseline; a positive, statistically significant excess at any order flags dependency and fails the Markov verdict.

Conclusions

Strict per-stream verdict and raw-hardware comparison

Stream	Bias	NIST	Next-bit	Markov	OVERALL
marrakesh_20260714-21080	0.0073	low	PASS	PASS	FAIL
marrakesh_20260714-210803	0.0002	OK	PASS	PASS	PASS
ed_marrakesh_20260714-1	0.0050	OK	PASS	PASS	PASS
ed_marrakesh_20260714-2	0.0031	OK	PASS	PASS	PASS

Hardware (raw) comparison: lower global bias = optimal_harvested_ibm_marrakesh_20260714-215818_processed (raw) (0.0050 vs 0.0073).

Strict next-bit = multi-window logistic-regression CV with Wilson CIs; a single predictable window fails the stream. Markov spans orders [1, 2, 4].

Note: whitened streams typically PASS everything — a good extractor whitens any high-entropy source — so judge the HARDWARE on the RAW rows, and read the whitened rows as 'what post-processing delivers'.

OVERALL PASS requires: NIST pass rate ≥ 0.90 , next-bit verdict PASS, and Markov verdict PASS. This is a guide, not a formal certification; sample size and calibration still matter.